

## Working in Linux

*Mocktar ISSA  
Full-Stack Software Engineer  
SBIN SA[CELTIS],*



# Compressing Files

**Compression** reduces the amount of data needed to store or transmit a file while storing it in such a way that the file can be restored. A file with human-readable text might have frequently used words replaced by something smaller, or an image with a solid background might represent patches of that color by a code. The compressed version of the file is not typically viewed or utilized, instead, it is decompressed before use.

**The compression algorithm** is a procedure the computer uses to encode the original file, and as a result, make it smaller. Computer scientists research these algorithms and come up with better ones that can work faster or make the input file smaller. When talking about compression, there are two types:

- **Lossless:** No information is removed from the file. Compressing a file and decompressing it leaves something identical to the original.
- **Lossy:** Information might be removed from the file. It is compressed in such a way that uncompressing a file will result in a file that is slightly different from the original. For instance, an image with two subtly different shades of green might be made smaller by treating those two shades as the same. Often, the eye can't pick out the difference anyway.

Linux provides several tools to compress files; the most common is **gzip**. Here we show a file before and after compression:

```
adavtyan@artur-lpt:~/tmp$ ls -l kern.log
-rw-r----- 1 adavtyan adavtyan 176788 26 14:55 kern.log
adavtyan@artur-lpt:~/tmp$ gzip kern.log
adavtyan@artur-lpt:~/tmp$ ls -l kern.log.gz
-rw-r----- 1 adavtyan adavtyan 28015 26 14:55 kern.log.gz
```

The **gzip** command will provide this information, by using the **-l** option, as shown here:

```
adavtyan@artur-lpt:~/tmp$ gzip -l kern.log.gz
      compressed      uncompressed      ratio uncompressed_name
      28015            176788      84.2% kern.log
```

Compressed files can be restored to their original form using either the **gunzip** command or the **gzip -d** command. This process is called decompression. After **gunzip** does its work, the kern.log file is restored to its original size and file name.

```
adavtyan@artur-lpt:~/tmp$ gunzip kern.log.gz
adavtyan@artur-lpt:~/tmp$ ls -l kern.log
-rw-r----- 1 adavtyan adavtyan 176788 26 14:55 kern.log
```

*Note: There are other commands that operate virtually identically to **gzip** and **gunzip**. There is **bzip2** and **bunzip2**, as well as **xz** and **unxz**.*

# Archiving Files

If you had several files to send to someone, you could choose to compress each one individually. You would have a smaller amount of data in total than if you sent uncompressed files, however, you would still have to deal with many files at one time.

*Archiving* is the solution to this problem. Archive files are typically used for a transfer (locally or over the internet) or make a backup copy of a collection of files and directories which allow you to work with only one file (if compressed, it has a lower size than the sum of all files within it) instead of many. Likewise, archives are used for software application packaging. This single file can be easily compressed for ease of transfer while the files in the archive retain the structure and permissions of the original files.

The traditional UNIX utility to archive files is called **tar**, which is a short form of TAPE Archive. It was used to stream many files to a tape for backups or file transfer. The **tar** command takes in several files and creates a single output file that can be split up again into the original files on the other end of the transmission.

The **tar** command has three modes that are helpful to become familiar with:

Mode	Function
<i>Create</i>	Make a new archive out of a series of files.
<i>Extract</i>	Pull one or more files out of an archive.
<i>List</i>	Show the contents of the archive without extracting.

Remembering the modes is key to figuring out the command line options necessary to do what you want. In addition to the mode, remember where to specify the name of the archive, as you may be entering multiple file names on a command line.

# Archiving Create Mode

To create an archive with tar, use the ‘-c’ (“create”) option, and specify the name of the archive file to create with the ‘-f’ option. It’s common practice to use a name with a ‘.tar’ extension. Note that unless specifically mentioned otherwise, all commands and command parameters used in the remainder of this article are used in lowercase.

```
tar -c [ -f ARCHIVE ] [ OPTION ] [ FILE ... ]
```

Creating an archive with the **tar** command requires two named options:

Option	Function
<b>-c</b>	Create an archive
<b>-f archive</b>	Use archive file. The argument <b>archive</b> will be the name of the resulting archive file

*Example:*

```
adavtyan@artur-lpt:~/tmp$ tar -cf kern.log.tar kern.log
adavtyan@artur-lpt:~/tmp$ ls -l kern.log.tar
-rw-rw-r-- 1 adavtyan adavtyan 184320 26 16:33 kern.log.tar
```

Normally, tarball files are slightly larger than the combined input files due to the overhead information on recreating the original files. Tarballs can be compressed for easier transport, either by using **gzip** on the archive or by having **tar** do it with the **-z** option.

```
adavtyan@artur-lpt:~/tmp$ tar -czf kern.log.tar.gz kern.log
adavtyan@artur-lpt:~/tmp$ ls -l kern.log.tar.gz
-rw-rw-r-- 1 adavtyan adavtyan 28134 26 16:37 kern.log.tar.gz
adavtyan@artur-lpt:~/tmp$ gzip -l kern.log.tar.gz
      compressed      uncompressed      ratio uncompressed_name
      28134           184320      84.7% kern.log.tar
```

# Archiving List Mode

To list the contents of a tar archive without extracting them, use tar with the ‘-t’ option.

```
tar -t [ -f ARCHIVE ] [ OPTION ]
```

The next example uses three options:

Option	Function
-t	List the files in an archive
-j	Decompress with bzip2 command
-f archive	Use archive file. The argument <b>archive</b> will be the name of the resulting archive file

*Example: Create new tar file*

```
adavtyan@artur-lpt:~/tmp$ tar -c -f kern.log.tar2 kern.log.tar.gz kern.log.tar kern.log
adavtyan@artur-lpt:~/tmp$ ls
kern.log  kern.log.tar  kern.log.tar2  kern.log.tar.gz
```

List the archive file

```
adavtyan@artur-lpt:~/tmp$ tar -t -f kern.log.tar2
kern.log.tar.gz
kern.log.tar
kern.log
adavtyan@artur-lpt:~/tmp$
```

*Note:* The **tar** command will recurse into subdirectories automatically when compressing and will store the path info inside the archive.



# Archiving Extract Mode

To extract (or *unpack*) the contents of a tar archive, use tar with the '-x' ("extract") option.

```
tar -x [ -f ARCHIVE ] [ OPTION ]
```

The next example uses three options:

Option	Function
-x	Extract files from an archive
-j	Decompress with bzip2 command
-f archive	Use archive file. The argument <b>archive</b> will be the name of the resulting archive file

Example:

```
adavtyan@artur-lpt:~/tmp/tar$ ls -la
total 220
drwxrwxr-x 2 adavtyan adavtyan  4096 UwJ  6 09:25 .
drwxrwxr-x 3 adavtyan adavtyan  4096 UwJ  6 09:25 ..
-rw-rw-r-- 1 adavtyan adavtyan 215040 UwJ  6 09:25 kern.log.tar
adavtyan@artur-lpt:~/tmp/tar$ tar -xf kern.log.tar
adavtyan@artur-lpt:~/tmp/tar$ ls -la
total 424
drwxrwxr-x 2 adavtyan adavtyan  4096 UwJ  6 09:26 .
drwxrwxr-x 3 adavtyan adavtyan  4096 UwJ  6 09:25 ..
-rw-r----- 1 adavtyan adavtyan 176788 UwJ 26 14:55 kern.log
-rw-rw-r-- 1 adavtyan adavtyan 215040 UwJ  6 09:25 kern.log.tar
-rw-rw-r-- 1 adavtyan adavtyan  28134 UwJ 26 16:37 kern.log.tar.gz
adavtyan@artur-lpt:~/tmp/tar$ tar -tf kern.log.tar
kern.log
kern.log.tar.gz
```

Note: Add the -v flag and you will get a verbose output of the files processed, making it easier to keep track of what's happening:

# Zip Files

The de facto archiving utility in Microsoft is the ZIP file. ZIP is not as prevalent in Linux but is well supported by the **zip** and **unzip** commands. Albeit, with **tar** and **gzip/gunzip** the same commands and options can be used interchangeably to do the creation and extraction, but this is not the case with **zip**. The same option has different meanings for the two different commands.

The default mode of **zip** is to add files to an archive and compress it.

```
zip [ OPTIONS ] [ zipfile [ file ... ] ]
```

The first argument zipfile is the name of the archive to be created, after that, a list of files to be added. The following example shows a compressed archive called files.zip being created:

```
adavtyan@artur-lpt:~/tmp/tar$ zip files.zip kern.log kafka.json cp-kafka-0.sh
  adding: kern.log (deflated 84%)
  adding: kafka.json (deflated 94%)
  adding: cp-kafka-0.sh (deflated 77%)
adavtyan@artur-lpt:~/tmp/tar$ ls -la
total 344
drwxrwxr-x 2 adavtyan adavtyan  4096 UwJ  6 09:32 .
drwxrwxr-x 3 adavtyan adavtyan  4096 UwJ  6 09:25 ..
-rw-rw-r-- 1 adavtyan adavtyan   2480 UwJ  6 09:32 cp-kafka-0.sh
-rw-rw-r-- 1 adavtyan adavtyan  35976 UwJ  6 09:32 files.zip
-rw-rw-r-- 1 adavtyan adavtyan 122485 UwJ  6 09:32 kafka.json
-rw-r----- 1 adavtyan adavtyan 176788 UwJ 26 14:55 kern.log
adavtyan@artur-lpt:~/tmp/tar$
```

It should be noted that **tar** requires the **-f** option to indicate a filename is being passed, while **zip** and **unzip** require a filename and therefore don't need you to inform the command a filename is being passed.

The **zip** command will not recurse into subdirectories by default, which is different behavior than the **tar** command. If you want **tar** like behavior, you must use the **-r** option to indicate recursion is to be used:

# Viewing Files in the Terminal

The **cat** command, short for concatenate, is a simple but useful command whose functions include creating and displaying text files, as well as combining copies of text files. One of the most popular uses of **cat** is to display the content of text files. To display a file in the standard output using the cat command, type the command followed by the filename:

```
adavtyan@artur-lpt:~/tmp/Documents$ cat food.txt
Food is good.
```

While viewing small files with the **cat** command poses no problems, it is not an ideal choice for large files. The **cat** command doesn't provide any easy ways to pause and restart the display, so the entire file contents are dumped to the screen.

For larger files, use a pager command to view the contents. Pager commands display one page of data at a time, allowing you to move forward and backward in the file by using movement keys.

There are two commonly used pager commands:

- The **less** command provides a very advanced paging capability. It is usually the default pager used by commands like the man command.
- The **more** command has been around since the early days of UNIX. While it has fewer features than the **less** command, however, the **less** command isn't included with all Linux distributions. The **more** command is always available.

The **more** and **less** commands allow users to move around the document using keystroke commands. Because developers based the **less** command on the functionality of the **more** command, all of the keystroke commands available in the **more** command also work in the **less** command.

The focus of our content is on the **more** advanced **less** command. The **more** command is still useful to remember for times when the **less** command isn't available. Remember that most of the keystroke commands provided work for both commands.



# Pager Movement Commands

To view a file with the **less** command, pass the file name as an argument:

```
less filename*
```

There are many movement commands for the **less** command, each with multiple possible keys or key combinations. While this may seem intimidating, it is not necessary to memorize all of these movement commands. When viewing a file with the **less** command, use the **H** key or **Shift+H** to display a help screen:

```
SUMMARY OF LESS COMMANDS

Commands marked with * may be preceded by a number, N.
Notes in parentheses indicate the behavior if N is given.
A key preceded by a caret indicates the Ctrl key; thus ^K is ctrl-K.

h H      Display this help.
q :q Q :Q ZZ  Exit.

-----

MOVING

e ^E j ^N CR * Forward one line (or N lines).
y ^Y k ^K ^P * Backward one line (or N lines).
f ^F ^V SPACE * Forward one window (or N lines).
b ^B ESC-V    * Backward one window (or N lines).
z          * Forward one window (and set window to N).
w          * Backward one window (and set window to N).
ESC-SPACE  * Forward one window, but don't stop at end-of-file.
d ^D      * Forward one half-window (and set half-window to N).
u ^U      * Backward one half-window (and set half-window to N).
```

There are two ways to search in the **less** command: searching forward or backward from your current position.

- To start a search to look forward from your current position, use the slash **/** key. Then, type the text or pattern to match and press the **Enter** key.
- To search backward from your current position, press the question mark **?** key, then type the text or pattern to match and press the **Enter** key. The cursor moves backward to the first match it can find or reports that the pattern cannot be found.

# Head and Tail

The **head** and **tail** commands are used to display only the first few or last few lines of a file, respectively (or, when used with a pipe, the output of a previous command). By default, the **head** and **tail** commands display ten lines of the file that is provided as an argument.

For example, the following command displays the first ten lines of the `/etc/sysctl.conf` file:

```
adavtyan@artur-lpt:~/tmp/Documents$ head /etc/sysctl.conf
#
# /etc/sysctl.conf - Configuration file for setting system variables
# See /etc/sysctl.d/ for additional system variables.
# See sysctl.conf (5) for information.
#
#kernel.domainname = example.com
# Uncomment the following to stop low-level messages on console
#kernel.printk = 3 4 1 3
```

Passing a number as an option will cause both the **head** and **tail** commands to output the specified number of lines, instead of the standard ten. For example to display the last five lines of the `/etc/sysctl.conf` file use the **-5** option:

```
adavtyan@artur-lpt:~/tmp/Documents$ tail -5 /etc/sysctl.conf
# 0=disable, 1=enable all, >1 bitmask of sysrq functions
# See https://www.kernel.org/doc/html/latest/admin-guide/sysrq.html
# for what other values do
#kernel.sysrq=438
adavtyan@artur-lpt:~/tmp/Documents$
```

Live file changes can be viewed by using the **-f** option to the **tail** command-useful when you want to see changes to a file as they are happening.

# Input/Output Redirection

Input/Output (I/O) redirection allows for command line information to be passed to different streams. Before discussing redirection, it is important to understand the standard streams.

Input and output in the Linux environment is distributed across three streams. These streams are:

1. **STDIN**  
Standard input, or STDIN, is information entered normally by the user via the keyboard. When a command prompts the shell for data, the shell provides the user with the ability to type commands that, in turn, are sent to the command as STDIN.
2. **STDOUT**  
Standard output, or STDOUT, is the normal output of commands. When a command functions correctly (without errors) the output it produces is called STDOUT. By default, STDOUT is displayed in the terminal window where the command is executing. STDOUT is also known as stream or channel #1.
3. **STDERR**  
Standard error, or STDERR, is error messages generated by commands. By default, STDERR is displayed in the terminal window where the command is executing. STDERR is also known as stream or channel #2.

The redirection capabilities built into Linux provide you with a robust set of tools used to make all sorts of tasks easier to accomplish. Whether you're writing complex software or performing file management through the command line, knowing how to manipulate the different I/O streams in your environment will greatly increase your productivity.

I/O redirection allows the user to redirect STDIN so that data comes from a file and STDOUT/STDERR so that output goes to a file. Redirection is achieved by using the arrow < > characters.

# STDOUT

**STDOUT** can be directed to files. To begin, observe the output of the following echo command which displays to the screen and Using the > character, the output can be redirected to a file instead:

```
adavtyan@artur-lpt:~/tmp/Documents$ echo "Line 1"
Line 1
adavtyan@artur-lpt:~/tmp/Documents$ echo "Line 1" > example.txt
adavtyan@artur-lpt:~/tmp/Documents$ ls
alpha-first.txt  example.txt  food.txt  kafka.json  new-home.txt  newhome.txt  profile.txt  red.txt  spelling.txt
adavtyan@artur-lpt:~/tmp/Documents$ cat example.txt
Line 1
```

It is important to realize that the single arrow overwrites any contents of an existing file:

```
adavtyan@artur-lpt:~/tmp/Documents$ echo "New line 1" > example.txt
adavtyan@artur-lpt:~/tmp/Documents$ cat example.txt
New line 1
```

The original contents of the file are gone, replaced with the output of the new echo command.

It is also possible to preserve the contents of an existing file by appending to it. Use two arrow >> characters to append to a file instead of overwriting it:

```
adavtyan@artur-lpt:~/tmp/Documents$ echo "Another line" >> example.txt
adavtyan@artur-lpt:~/tmp/Documents$ cat example.txt
New line 1
Another line
```

Instead of being overwritten, the output of the echo command is added to the bottom of the file.

# STDERR

**STDERR** can be redirected similarly to STDOUT. When using the arrow character to redirect, stream #1 (STDOUT) is assumed unless another stream is specified. Thus, stream #2 must be specified when redirecting STDERR by placing the number 2 preceding the arrow > character.

To demonstrate redirecting STDERR, first observe the following command which produces an error because the specified directory does not exist:

```
adavtyan@artur-lpt:~/tmp/Documents$ ls /fake
ls: cannot access '/fake': No such file or directory
```

Note that there is nothing in the example above that implies that the output is STDERR. The output is clearly an error message, but how could you tell that it is being sent to STDERR? One easy way to determine this is to redirect STDOUT:

```
adavtyan@artur-lpt:~/tmp/Documents$ ls /fake > output.txt
ls: cannot access '/fake': No such file or directory
```

In the example above, STDOUT was redirected to the output.txt file. So, the output that is displayed can't be STDOUT because it would have been placed in the output.txt file instead of the terminal. Because all command output goes either to STDOUT or STDERR, the output displayed above must be STDERR.

The STDERR output of a command can be sent to a file:

```
adavtyan@artur-lpt:~/tmp/Documents$ ls /fake 2> error.txt
adavtyan@artur-lpt:~/tmp/Documents$ cat error.txt
ls: cannot access '/fake': No such file or directory
```



# Redirecting Multiple Streams

It is possible to direct both the STDOUT and STDERR of a command at the same time. The following command produces both STDOUT and STDERR because one of the specified directories exists and the other does not:

```
adavtyan@artur-lpt:~/tmp/Documents$ ls /fake /etc/ppp
ls: cannot access '/fake': No such file or directory
/etc/ppp:
chap-secrets  ip-down  ip-down.d  ip-up  ip-up.d  ipv6-down  ipv6-down.d  ipv6-up  ipv6-up.d  options  options.pptp  pap-secrets  peers  resolv.conf
```

If only the STDOUT is sent to a file, STDERR is still printed to the screen and If only the STDERR is sent to a file, STDOUT is still printed to the screen:

```
adavtyan@artur-lpt:~/tmp/Documents$ ls /fake /etc/ppp > example.txt
ls: cannot access '/fake': No such file or directory
adavtyan@artur-lpt:~/tmp/Documents$ ls /fake /etc/ppp 2> error.txt
/etc/ppp:
chap-secrets  ip-down  ip-down.d  ip-up  ip-up.d  ipv6-down  ipv6-down.d  ipv6-up  ipv6-up.d  options  options.pptp  pap-secrets  peers  resolv.conf
adavtyan@artur-lpt:~/tmp/Documents$ cat error.txt
ls: cannot access '/fake': No such file or directory
adavtyan@artur-lpt:~/tmp/Documents$ cat example.txt
/etc/ppp:
chap-secrets
ip-down
ip-down.d
```

Both STDOUT and STDERR can be sent to a file by using the ampersand & character in front of the arrow > character. The &> character set means both 1> and 2> or If you don't want STDERR and STDOUT to both go to the same file, they can be redirected to different files by using both > and 2>. For example, to direct STDOUT to example.txt and STDERR to error.txt execute the following:

```
adavtyan@artur-lpt:~/tmp/Documents$ ls /fake /etc/ppp /junk /etc/sound &> all.txt
adavtyan@artur-lpt:~/tmp/Documents$ ls /fake /etc/ppp > example.txt 2> error.txt
adavtyan@artur-lpt:~/tmp/Documents$ █
```

# STDINN

The concept of redirecting STDIN is a difficult one because it is more difficult to understand why you would want to redirect STDIN. With STDOUT and STDERR, their purpose is straightforward; sometimes it is helpful to store the output into a file for future use.

Most Linux users end up redirecting STDOUT routinely, STDERR on occasion, and STDIN very rarely.

There are very few commands that require you to redirect STDIN because with most commands if you want to read data from a file into a command, you can specify the filename as an argument to the command.

For some commands, if you don't specify a filename as an argument, they revert to using STDIN to get data. For example, consider the following cat command:

```
adavtyan@artur-lpt:~/tmp/Documents$ cat
hello
hello
how are you ?
how are you ?
goodbye
goodbye
^C
```

The first command in the example below redirects the output of the cat command to a newly created file called new.txt. This action is followed up by providing the cat command with the *new.txt* file as an argument to display the redirected text in STDOUT.

```
adavtyan@artur-lpt:~/tmp/Documents$ cat > new.txt
Hello
Bye
^C
adavtyan@artur-lpt:~/tmp/Documents$ cat new.txt
Hello
Bye
```

# Sorting files or Input

The sort command can be used to rearrange the lines of files or input in either dictionary or numeric order. The following example creates a small file, using the head command to grab the first 3 lines of the /etc/passwd file and send the output to a file called mypasswd.

```
adavtyan@artur-lpt:~/tmp$ head -n 3 /etc/passwd > mypasswd
adavtyan@artur-lpt:~/tmp$ cat mypasswd
root:x:0:0:root:/root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin
bin:x:2:2:bin:/bin:/usr/sbin/nologin
adavtyan@artur-lpt:~/tmp$
```

Now we will **sort** the mypasswd file:

```
adavtyan@artur-lpt:~/tmp$ sort mypasswd
bin:x:2:2:bin:/bin:/usr/sbin/nologin
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin
root:x:0:0:root:/root:/bin/bash
adavtyan@artur-lpt:~/tmp$
```

Three options are used to achieve:

Option	Function
-t	The -t option specifies the field delimiter. If the file or input is separated by a delimiter other than whitespace, for example a comma or colon, the -t option will allow for another field separator to be specified as an argument.
-k	The -k option specifies the field number. To specify which field to sort by, use the -k option with an argument to indicate the field number, starting with 1 for the first field.
-n	This option specifies the sort type.

Example:

```
adavtyan@artur-lpt:~/tmp$ sort -t: -n -k3 mypasswd
root:x:0:0:root:/root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin
bin:x:2:2:bin:/bin:/usr/sbin/nologin
adavtyan@artur-lpt:~/tmp$
```

## Filter File Contents [grep]

The grep command can be used to filter lines in a file or the output of another command that matches a specified pattern. That pattern can be as simple as the exact text that you want to match or it can be much more advanced through the use of regular expressions.

For example, to find all the users who can log in to the system with the BASH shell, the grep command can be used to filter the lines from the /etc/passwd file for the lines containing the pattern bash:

```
adavtyan@artur-lpt:~/tmp$ grep bash /etc/passwd
root:x:0:0:root:/root:/bin/bash
adavtyan:x:1000:1000:Artur Davtyan,,,:/home/adavtyan:/bin/bash
```

In some cases, it may not be important to find the specific lines that match the pattern, but rather how many lines match the pattern. The -c option provides a count of how many lines match:

```
adavtyan@artur-lpt:~/tmp$ grep -c bash /etc/passwd
2
```

The -n option to the grep command will display original line numbers. To display all lines and their line numbers in the /etc/passwd file which contain the pattern bash:

```
adavtyan@artur-lpt:~/tmp$ grep -n bash /etc/passwd
1:root:x:0:0:root:/root:/bin/bash
41:adavtyan:x:1000:1000:Artur Davtyan,,,:/home/adavtyan:/bin/bash
```

The -v option inverts the match, outputting all lines that do not contain the pattern. To display all lines not containing nologin in the /etc/passwd file:

```
adavtyan@artur-lpt:~/tmp$ grep -v nologin /etc/passwd
root:x:0:0:root:/root:/bin/bash
sync:x:4:65534:sync:/bin:/bin/sync
speech-dispatcher:x:111:29:Speech Dispatcher,,,:/var/run/speech-dispatcher:/bin/false
whoopsie:x:112:117:/:nonexistent:/bin/false
```

## Filter File Contents [grep]

The **-i** option ignores the case (capitalization) distinctions. The following searches for the pattern **the** in **newhome.txt**, allowing each character to be uppercase or lowercase:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep the newhome.txt
**Beware** of the ghost in the bedroom.
**Caution** the spirits don't like guests.
adavtyan@artur-lpt:~/tmp/Documents$ grep -i the newhome.txt
There are three bathrooms.
**Beware** of the ghost in the bedroom.
The kitchen is open for entertaining.
**Caution** the spirits don't like guests.
```

The **-w** option only returns lines which contain matches that form whole words. To be a word, the character string must be preceded and followed by a non-word character. Word characters include letters, digits, and the underscore character.

The following examples search for the **are** pattern in the **newhome.txt** file. The first command searches with no options, while the second command includes the **-w** option. Compare the outputs:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep are newhome.txt
There are three bathrooms.
**Beware** of the ghost in the bedroom.
adavtyan@artur-lpt:~/tmp/Documents$ grep -w are newhome.txt
There are three bathrooms.
adavtyan@artur-lpt:~/tmp/Documents$ grep -nw are newhome.txt
1:There are three bathrooms.
adavtyan@artur-lpt:~/tmp/Documents$
```



# Basic Regular Expressions

Regular expressions, also referred to as regex, are a collection of normal and special characters that are used to find simple or complex patterns, respectively, in files. These characters are characters that are used to perform a particular matching function in a search.

Normal characters are alphanumeric characters which match themselves. For example, an a would match an a. Special characters have special meanings when used within patterns by commands like the **grep** command. They behave in a more complex manner and do not match themselves.

There are both Basic Regular Expressions (available to a wide variety of Linux commands) and Extended Regular Expressions (available to more advanced Linux commands). Basic Regular Expressions include the following:

Character	Function
.	Any single character
[ ]	A list or range of characters to match one character If the first character within the brackets is the caret ^, it means any character not in the list
*	The previous character repeated zero or more times
^	If the first character in the pattern, the pattern must be at the beginning of the line to match, otherwise just a literal ^ character
\$	If the last character in the pattern, the pattern must be at the end of the line to match, otherwise just a literal \$ character

The **grep** command is just one of the many commands that support regular expressions. Some other commands include the more and less commands.

*Note: While some of the regular expressions are unnecessarily quoted with single quotes, it is good practice to use single quotes around regular expressions to prevent the shell from trying to interpret special meaning from them.*

## The Period . Character

One of the most useful expressions is the period . character. It matches any character except for the new line character. Consider the unfiltered contents of the ~/Documents/red.txt file:

```
adavtyan@artur-lpt:~/tmp/Documents$ cat red.txt
red
reef
rot
reeed
rd
rod
roof
reed
root
reel
read
adavtyan@artur-lpt:~/tmp/Documents$ grep 'r..f' red.txt
reef
roof
```

The line does not have to be an exact match, it simply must contain the pattern, as seen here when r..t is searched for in the /etc/passwd file:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep 'r..t' /etc/passwd
root:x:0:0:root:/root:/bin/bash
nm-openvpn:x:126:134:NetworkManager OpenVPN,,,:/var/lib/openvpn/chroot:/usr/sbin/nologin
adavtyan@artur-lpt:~/tmp/Documents$
```

The period character can be used any number of times. To find all words that have at least four characters, the following pattern can be used:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep '....' red.txt
reef
reeed
roof
reed
root
reel
read
```

## The Bracket [] Character

When using the `.` character, any possible character could match it. In some cases, you want to specify exactly which characters you want to match, such as a lowercase alphabet character or a number character.

The square brackets `[]` match a single character from the list or range of possible characters contained within the brackets. To find all the lines in `profile.txt` which have a number in them, use the pattern `[0123456789]` or `[0-9]`:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep '[0-9]' profile.txt
I am 37 years old.
3121991
I have 2 dogs.
123456789101112
adavtyan@artur-lpt:~/tmp/Documents$
```

Note that each possible character can be listed out `[abcd]` or provided as a range `[a-d]`, as long as the range is in the correct order. For example, `[d-a]` wouldn't work because it isn't a valid range:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep '[d-a]' profile.txt
grep: Invalid range end
adavtyan@artur-lpt:~/tmp/Documents$
```

To match a character that is not one of the listed characters, start the set with a `^` symbol. To find all the lines which contain any non-numeric characters, insert a `^` as the first character inside the brackets. This character negates the characters listed:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep '[^0-9]' profile.txt
Hello my name is Joe.
I am 37 years old.
My favorite food is avocados.
I have 2 dogs.
adavtyan@artur-lpt:~/tmp/Documents$
```

# The Asterisk \* Character

The asterisk \* character is used to match zero or more occurrences of a character or pattern preceding it. For example, e\* would match zero or more occurrences of the letter e:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep 're*d' red.txt
red
reeed
rd
reed
```

It is also possible to match zero or more occurrences of a list of characters by utilizing the square brackets. The pattern [oe]\* used in the following example matches zero or more occurrences of the o character or the e character:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep 'r[oe]*d' red.txt
red
reeed
rd
rod
reed
```

When used with only one other character, \* isn't very helpful. Any of the following patterns would match every string or line in the file: '.\*' 'e\*' 'b\*' 'z\*' because the asterisk \* character can match zero occurrences of a pattern.

```
adavtyan@artur-lpt:~/tmp/Documents$ grep 'z*' red.txt
red
reef
rot
reeed
```

To make the asterisk character useful, it is necessary to create a pattern which includes more than just the one character preceding it. For example, the results above can be refined by adding another e to make the pattern ee\* effectively matching every line which contains at least one e.

```
adavtyan@artur-lpt:~/tmp/Documents$ grep 'ee*' red.txt
red
reef
reeed
reed
reel
read
```

# Anchor Character

When performing a pattern match, the match could occur anywhere on the line. Anchor characters are one of the ways regular expressions can be used to narrow down search results. They specify whether the match occurs at the beginning of the line or the end of the line.

For example, the pattern `root` appears many times in the `/etc/passwd` file:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep 'root' /etc/passwd
root:x:0:0:root:/root:/bin/bash
nm-openvpn:x:126:134:NetworkManager OpenVPN,,,:/var/lib/openvpn/chroot:/usr/sbin/nologin
adavtyan@artur-lpt:~/tmp/Documents$
```

The caret (circumflex) `^` character is used to ensure that a pattern appears at the beginning of the line. For example, to find all lines in `/etc/passwd` that start with `root` use the pattern `^root`. Note that `^` must be the first character in the pattern to be effective:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep '^root' /etc/passwd
root:x:0:0:root:/root:/bin/bash
adavtyan@artur-lpt:~/tmp/Documents$
```

The second anchor character `$` can be used to ensure a pattern appears at the end of the line, thereby effectively reducing the search results. To find the lines that end with an `r` in the `alpha-first.txt` file, use the pattern `r$`:

```
adavtyan@artur-lpt:~/tmp/Documents$ cat alpha-first.txt
A is for Animal
B is for Bear
C is for Cat
D is for Dog
E is for Elephant
F is for Flower
adavtyan@artur-lpt:~/tmp/Documents$ grep 'r$' alpha-first.txt
B is for Bear
F is for Flower
adavtyan@artur-lpt:~/tmp/Documents$
```

Again, the position of this character is important. The `$` must be the last character in the pattern to be effective as an anchor.



# The Backslash \ Character

In some cases, you may want to match a character that happens to be a special regular expression character. For example, consider the following:

```
adavtyan@artur-lpt:~/tmp/Documents$ cat new-home.txt
Thanks for purchasing your new home!!

**Warning** it may be haunted.

There are three bathrooms.

**Beware** of the ghost in the bedroom.

The kitchen is open for entertaining.

**Caution** the spirits don't like guests.

Good luck!!!
adavtyan@artur-lpt:~/tmp/Documents$ grep 're*' new-home.txt
Thanks for purchasing your new home!!
**Warning** it may be haunted.
There are three bathrooms.
**Beware** of the ghost in the bedroom.
The kitchen is open for entertaining.
**Caution** the spirits don't like guests.
```

In the output of the grep command above, the search for `re*` matched every line which contained an `r` followed by zero or more of the letter `e`. To look for an actual asterisk `*` character, place a backslash `\` character before the asterisk `*` character:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep 're\*' new-home.txt
**Beware** of the ghost in the bedroom.
adavtyan@artur-lpt:~/tmp/Documents$
```

# Extended Regular Expressions

The use of extended regular expressions often requires a special option be provided to the command to recognize them. Historically, there is a command called `egrep`, which is similar to `grep`, but can understand extended regular expressions. Now, the `egrep` command is deprecated in favor of using `grep` with the `-E` option.

The following regular expressions are considered extended:

Character	Function
<code>?</code>	Matches previous character zero or one time, so it is an optional character
<code>+</code>	Matches previous character repeated one or more times
<code> </code>	Alternation or like a logical "or" operator

To match “colo” followed by zero or one `u` character followed by an `r` character:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep -E 'colou?r' spelling.txt
American English: Do you consider gray to be a color or a shade?
British English: Do you consider grey to be a colour or a shade?
```

To match one or more `e` characters:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep -E 'e+' red.txt
red
reef
reeed
reed
reel
read
```

To match either `gray` or `grey`:

```
adavtyan@artur-lpt:~/tmp/Documents$ grep -E 'gray|grey' spelling.txt
American English: Do you consider gray to be a color or a shade?
British English: Do you consider grey to be a colour or a shade?
```

## Memory usage check

The motherboard typically has slots where random-access memory (RAM) can be connected to the system. The 32-bit architecture systems can use up to 4 gigabytes (GB) of RAM, while 64-bit architectures are capable of addressing and using far more RAM.

In some cases, the RAM a system has might not be enough to handle all of the operating system requirements. Once invoked, programs are loaded in RAM along with any data they need to store, while instructions are sent to the processor when they execute.

To view the amount of RAM in your system, including the swap space, execute the **free** command. The **free** command has a **-m** option to force the output to be rounded to the nearest megabyte (MB) and a **-g** option to force the output to be rounded to the nearest gigabyte (GB):

```
adavtyan@artur-lpt:~/tmp/Documents$ free -m
```

	total	used	free	shared	buff/cache	available
Mem:	31911	4816	22301	332	4793	26307
Swap:	979	0	979			

If you want to monitor memory usage over time with the **free** command, then you can execute it with the **-s** option (how often to update) and specify that number of seconds. For example, executing the following **free** command would update the output every ten seconds:

```
adavtyan@artur-lpt:~/tmp/Documents$ free -m -s 20
```

	total	used	free	shared	buff/cache	available
Mem:	31911	5423	20729	648	5758	25385
Swap:	979	0	979			

  

	total	used	free	shared	buff/cache	available
Mem:	31911	5376	20776	648	5758	25431
Swap:	979	0	979			

# Processes

The kernel provides access to information about active processes through a pseudo filesystem that is visible under the `/proc` directory. Hardware devices are made available through special files under the `/dev` directory, while information about those devices can be found in another pseudo filesystem under the `/sys` directory.

Pseudo *filesystems* appear to be real files on disk but exist only in memory. Most pseudo file systems such as `/proc` are designed to appear to be a hierarchical tree off the root of the system of directories, files and subdirectories, but in reality only exist in the system's memory, and only appear to be resident on the storage device that the root file system is on.

The `/proc` directory not only contains information about running processes, as its name would suggest, but it also contains information about the system hardware and the current kernel configuration.

The `/proc` directory is read, and its information utilized by many different commands on the system, including but not limited to `top`, `free`, `mount`, `umount` and many many others. It is rarely necessary for a user to mine the `/proc` directory directly—it's easier to use the commands that utilize its information.

```
adavtyan@artur-lpt:~/tmp/Documents$ ls /proc
1      1060 1180 13903 147 159 17654 2157 2232 2444 2593 2706 2914 2989 3041 33 34821 36724 39 4662 53 646 743 89 dma locks sys
10     1065 1182 13905 1472 16 18 2165 2242 2459 26 2708 29247 2990 3056 3306 34848 36886 4 4675 54 65 744 9 driver mdstat sysrq-trigger
1008   1070 1192 13925 148 160 180 2170 2246 2476 2611 273 2926 2991 30627 332 34849 36933 40 4691 541 650 745 90 execdomains meminfo sysvipc
1032   1071 12 13927 1489 16061 18002 2171 2252 2478 2613 274 2930 2992 308 3326 35 36941 41 47 545 654 75 91 fb misc thread-self
1033   1072 1200 13940 149 161 189 2177 2256 2479 2618 275 2935 2995 309 33272 35013 37020 4151 4736 546 66 751 917 filesystems modules timer_list
1034   1073 1202 13941 15 162 1911 2179 2262 2480 2620 2765 2943 29953 30920 33333 35090 37027 42 4771 553 662 76 934 fs mounts tty
1040   1074 1206 14 1504 163 1913 2180 23 2481 2625 278 2951 29972 3093 3346 35152 37111 44 48 554 667 77 935 interrupts mtrr uptime
1041   1075 1208 140 151 164 1915 2182 2375 2482 2627 28 29539 2999 3096 33521 35216 37123 45 485 56 668 78 982 iomem net version
1042   1076 1211 141 153 16416 192 2185 2376 2483 2644 2832 29613 3 3097 33538 35441 37125 4544 486 57 68 80 acpi ioports pagetypeinfo version_signature
1043   11 1215 14450 154 16424 1923 2186 2377 2484 2649 2849 29639 30 310 33599 35456 37336 4545 4917 58 688 81 asound irq partitions vmallocinfo
1046   1105 1216 14453 15544 1644 1926 2193 2380 2485 2651 2854 2965 3002 311 3363 35649 37406 46 493 59 69 82 buddyinfo pressure vmstat
1048   1106 1217 145 156 165 1927 22 2385 2486 2665 2860 2968 3003 31249 3366 35736 37532 4617 494 6 70 83 bus kcore sched_debug zoneinfo
1049   1156 1225 14527 1560 1657 1929 2208 24 2487 2669 2867 2975 3004 3146 33670 36 37806 462 495 60 705 84 cgroups keys schedstat
1050   1157 1295 14536 157 17 1964 22118 2406 2489 2671 2881 2978 3005 3152 33681 36193 37809 4624 498 609 71 856 cmdline key-users scsi
1052   1160 13 14556 1572 17091 2 2214 2422 2492 2678 29 2979 30097 3163 337 36370 37897 4625 50 62 72 86 consoles kmsg self
1053   1162 1319 14557 15759 17142 20 2220 2425 2495 2683 2904 2980 302 31909 3376 36467 37937 4629 500 63 735 866 cpuinfo kpagecgroup slabinfo
1056   1165 1352 146 15763 172 207 2225 2428 2503 2688 2908 2982 30205 32 338 36514 37938 4631 5099 637 738 87 crypto kpagecount softirqs
1057   1169 139 14638 158 173 21 2230 2431 2546 2692 2909 2985 3029 32549 34 36630 38 4633 51 64 74 88 devices kpageflags stat
1059   1173 13901 1464 15818 17366 2154 22306 2440 2547 27 2911 2988 304 32684 3405 36632 38144 4636 52 642 742 881 diskstats loadavg swaps
```

The output shows a variety of named and numbered directories. There is a numbered directory for each running process on the system, where the name of the directory matches the process ID (PID) for the running process.

# Process Hierarchy

When the kernel finishes loading during the boot procedure, it starts the init process and assigns it a PID of 1. This process then starts other system processes, and each process is assigned a PID in sequential order.

As either of the init processes starts up other processes, they, in turn, may start up processes, which may start up other processes, on and on. When one process starts another process, the process that performs the starting is called the parent process and the process that is started is called the child process. When viewing processes, the parent PID is labeled PPID.

When the system has been running for a long time, it may eventually reach the maximum PID value, which can be viewed and configured through the `/proc/sys/kernel/pid_max` file. Once the largest PID has been used, the system "rolls over" and continues seamlessly by assigning PID values that are available at the bottom of the range.

Processes can be "mapped" into a family tree of parent and child couplings. If you want to view this tree, the command `ps tree` displays it:

```
adavtyan@artur-lpt:~/tmp/Documents$ ps tree
systemd--ModemManager--2*[{ModemManager}]
--NetworkManager--2*[{NetworkManager}]
--accounts-daemon--2*[{accounts-daemon}]
--acpid
--apache2--2*[apache2--26*[{apache2}]]
--at-spi-bus-laun--dbus-daemon
--at-spi2-registr--2*[{at-spi2-registr}]
--avahi-daemon--avahi-daemon
--bdepsecd--30*[{bdepsecd}]
--bdlogd--11*[{bdlogd}]
--bdmond--3*[{bdmond}]
--bdregd--2*[{bdregd}]
--bdsrvscand.bin--95*[{bdsrvscand.bin}]
--bdupdated--4*[{bdupdated}]
```



## Viewing Process Snapshot

Another way of viewing processes is with the `ps` command. By default, the `ps` command only shows the current processes running in the current shell. Ironically, even though you are trying to obtain information about processes, the `ps` command includes itself in the output:

```
adavtyan@artur-lpt:~/tmp/Documents$ ps
  PID TTY          TIME CMD
 16424 pts/1    00:00:01 bash
 35244 pts/1    00:00:00 ps
```

If you run `ps` with the option `--forest`, then, similar to the `pstree` command, it shows lines indicating the parent and child relationship:

```
adavtyan@artur-lpt:~/tmp/Documents$ ps --forest
  PID TTY          TIME CMD
 16424 pts/1    00:00:01 bash
 35425 pts/1    00:00:00 \_ ps
```

To be able to view all processes on the system execute either the `ps aux` command or the `ps -ef` command:

```
adavtyan@artur-lpt:~/tmp/Documents$ ps -ef | head -n 4
UID          PID    PPID  C STIME TTY          TIME CMD
root           1         0  0 00:00 ?        00:00:05 /sbin/init splash
root           2         0  0 00:00 ?        00:00:00 [kthreadd]
root           3         2  0 00:00 ?        00:00:00 [rcu_gp]
```

An administrator may be more concerned about the processes of another user. There are several styles of options that the `ps` command supports, resulting in different ways to view an individual user's processes. To use the traditional UNIX option to view the processes of a specific user, use the `-u` option:

```
adavtyan@artur-lpt:~/tmp/Documents$ ps -u root
  PID TTY          TIME CMD
    1 ?            00:00:05 systemd
    2 ?            00:00:00 kthreadd
    3 ?            00:00:00 rcu_gp
```

# Viewing Process in Real Time

Whereas the `ps` command provides a snapshot of the processes running at the instant the command is executed, the `top` command has a dynamic, screen-based interface that regularly updates the output of running processes.

By default, the output of the `top` command is sorted by the percentage % of CPU time that each process is currently using, with the higher values listed first, meaning more CPU-intensive processes are listed first:

```
top - 15:04:51 up 20:09, 1 user, load average: 0.89, 0.56, 0.39
Tasks: 452 total, 2 running, 450 sleeping, 0 stopped, 0 zombie
%Cpu(s): 5.9 us, 1.9 sy, 0.0 ni, 91.3 id, 0.1 wa, 0.0 hi, 0.7 si, 0.0 st
MiB Mem : 31911.5 total, 20847.7 free, 5350.6 used, 5713.2 buff/cache
MiB Swap: 980.0 total, 980.0 free, 0.0 used, 25498.4 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
2881	adavtyan	20	0	4739944	587732	124444	R	24.2	1.8	24:58.24	gnome-shell
14527	adavtyan	20	0	32.5g	226416	106496	S	22.8	0.7	18:43.43	chrome
36772	adavtyan	20	0	36.4g	103296	79448	S	10.9	0.3	0:00.33	chrome
29953	adavtyan	20	0	36.6g	287156	234396	S	10.6	0.9	2:03.93	chrome
2230	root	-51	0	0	0	0	S	8.6	0.0	8:39.87	irq/147-nvidia

There is an extensive amount of interactive commands that can be executed from within the running `top` program. Use the **H** key to view a full list.

One of the advantages of the `top` command is that it can be left running to stay on top of processes for monitoring purposes. If a process begins to dominate, or run away with the system, then by default it will appear at the top of the list presented by the `top` command.

```
adavtyan@artur-lpt:~/tmp/Documents$ cat /proc/loadavg
0.58 0.57 0.42 1/1651 37000
```

- **Load Average:** The first three numbers in this file indicate the load average over the last one, five and fifteen minute intervals.
- **Number of processes:** The fourth value is a fraction which shows the number of processes currently executing code on the CPU 1 and the total number of processes 1651.
- **Last PID :** The fifth value is the last PID value that executed code on the CPU.

# Log Files

As the kernel and various processes run on the system, they produce output that describes how they are running. Some of this output is displayed as standard output and error in the terminal window where the process was executed, though some of this data is not sent to the screen. Instead, it is written to various files. This information is called log data or log messages.

Log files are useful for many reasons; they help troubleshoot problems and determine whether or not unauthorized access has been attempted.

Some processes can log their own data to these files, other processes rely on a separate process (a daemon) to handle these log data files.

In yet more recent distributions, those based on systemd, the logging daemon is named `journald`, and the logs are designed to allow for mainly text output, but also binary. The standard method for viewing `journald`-based logs is to use the `journalctl` command.

You can view the contents of various log files using two different methods. First, as with most other files, you can use the `cat` command, or the `less` command to allow for searching, scrolling and other options.

The second method is to use the `journalctl` command on systemd-based systems, mainly because the `/var/log/journal` file now often contains binary information and using the `cat` or `less` commands may produce confusing screen behavior from control codes and binary items in the log files.

Although most log files contain text as their contents, which can be viewed safely with many tools, other files such as the `/var/log/btmp` and `/var/log/wtmp` files contain binary. By using the `file` command, users can check the file content type before they view it to make sure that it is safe to view. The following `file` command classifies `/var/log/wtmp` as data, which usually means the file is binary:

```
adavtyan@artur-lpt:~/tmp/Documents$ file /var/log/wtmp
/var/log/wtmp: dBase III DBT, version number 0, next free block index 2
```

# Log Files

Regardless of what the daemon process being used, the log files themselves are almost always placed into the `/var/log` directory structure. Although some of the file names may vary, here are some of the more common files to be found in this directory:

File	Contents
<code>boot.log</code>	Messages generated as services are started during the startup of the system.
<code>cron</code>	Messages generated by the <code>crond</code> daemon for jobs to be executed on a recurring basis.
<code>dmesg</code>	Messages generated by the kernel during system boot up.
<code>maillog</code>	Messages produced by the mail daemon for e-mail messages sent or received.
<code>messages</code>	Messages from the kernel and other processes that don't belong elsewhere. Sometimes named <code>syslog</code> instead of <code>messages</code> after the daemon that writes this file.
<code>secure</code>	Messages from processes that required authorization or authentication (such as the login process).
<code>journal</code>	Messages from the default configuration of the <code>systemd-journald.service</code> ; can be configured in the <code>/etc/journald.conf</code> file amongst other places.
<code>Xorg.0.log</code>	Messages from the X Windows (GUI) server.

**Thank you for your attention !**

**Q&A**

